

Cloud Programming: Lecture6 – Spark

***National Tsing-Hua University
2016, Spring Semester***



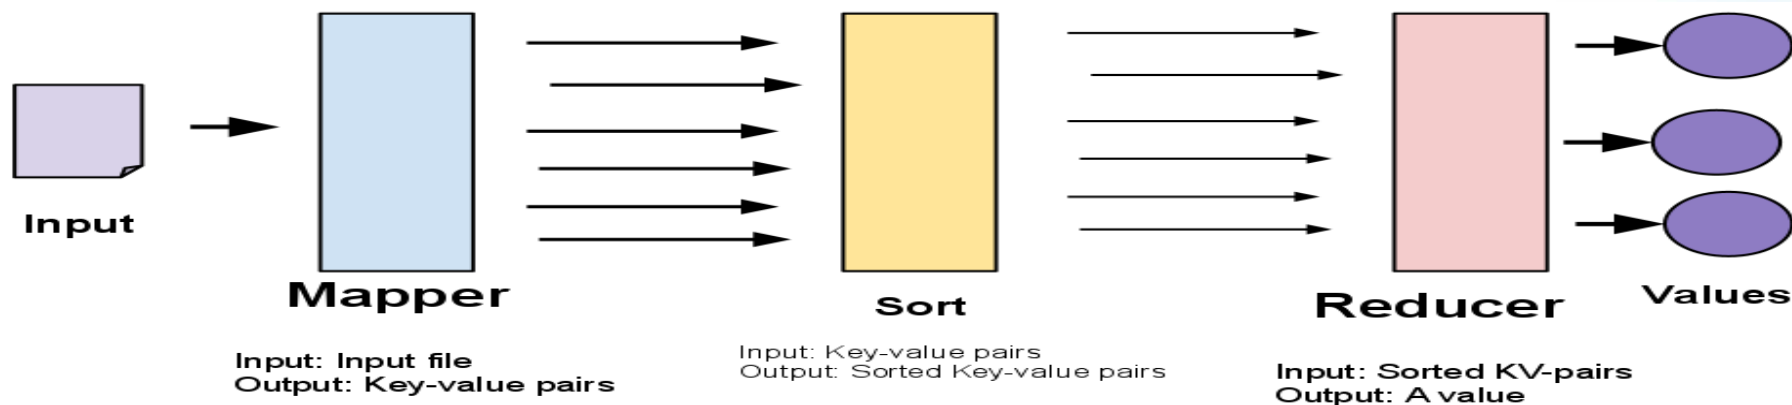
國立清華大學

Outline

- Overview of Hadoop/MapReduce Limitation
- Spark: In-memory computing

Limitation of MapReduce

- Simple but limited programming model



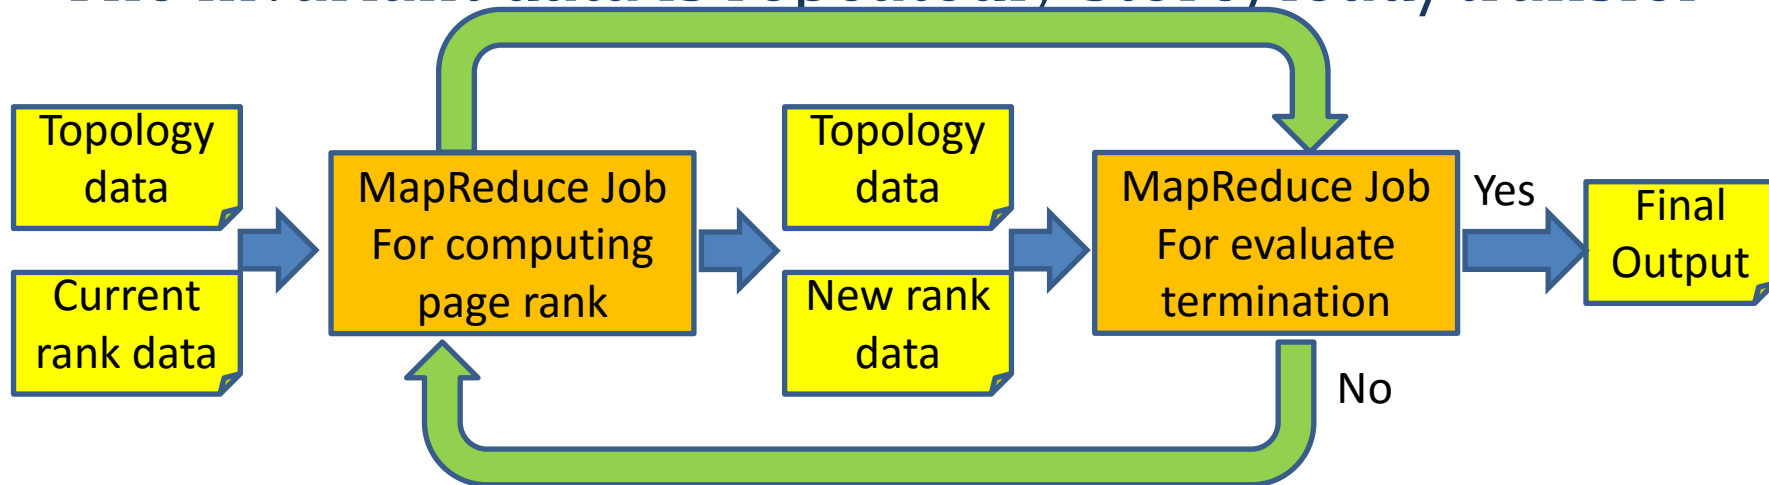
- Can only apply two computation functions in a job: Map&Reduce
 - ➔ More complex work must use **multiple** jobs
- The input and output of a job must store into a FS
 - ➔ FS(disk) is the only device to provide **data persistency**

Iterative Data Processing

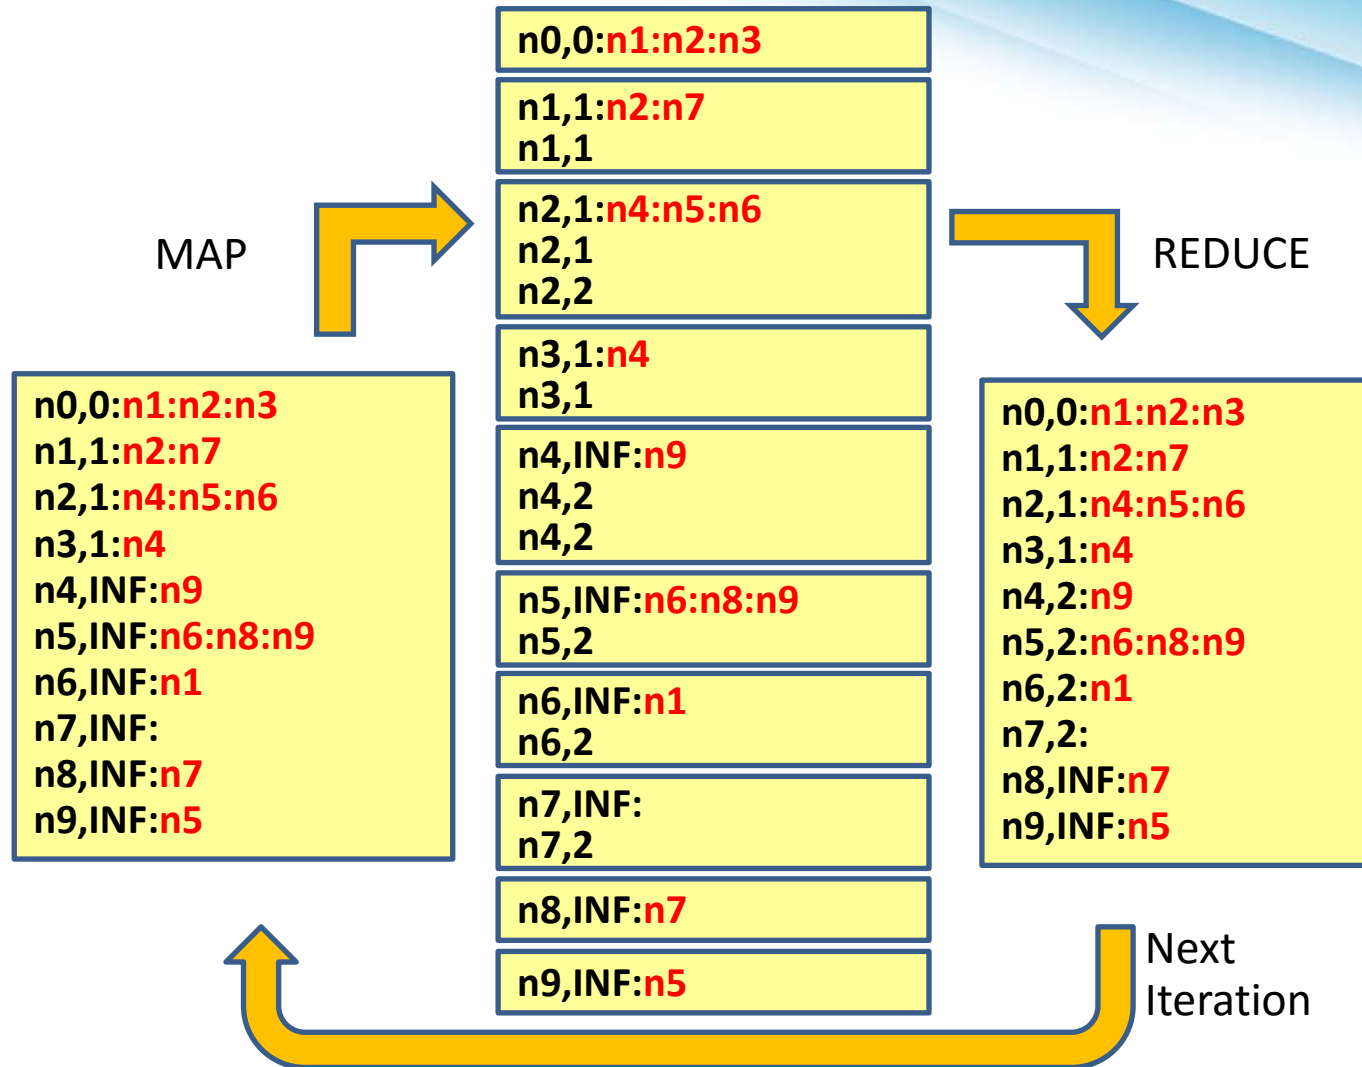
- **Definition:**
 - A mathematical procedure that generates a sequence of improving **approximate solutions** for a class of problems
 - The procedure iterates until converge or reach some termination criteria
- **Property of computation**
 - Termination criteria must be evaluated after each iteration
 - The output from an iteration will be the input of the next iteration
 - Invariant data must re-load and re-processing in the loop
- **Applications:**
 - Machine learning algorithm: K-mean
 - Graph algorithm: Page-rank
 - Approximation algorithm

Iterative Method in MapReduce

- Each iteration is submitted as an independent job
 - hard to ask the scheduler to manage and guarantee the performance of the whole application
- Termination criteria is evaluated after each iteration
- Data is written and read out disk after each iteration
- The invariant data is repeatedly store/load/transfer

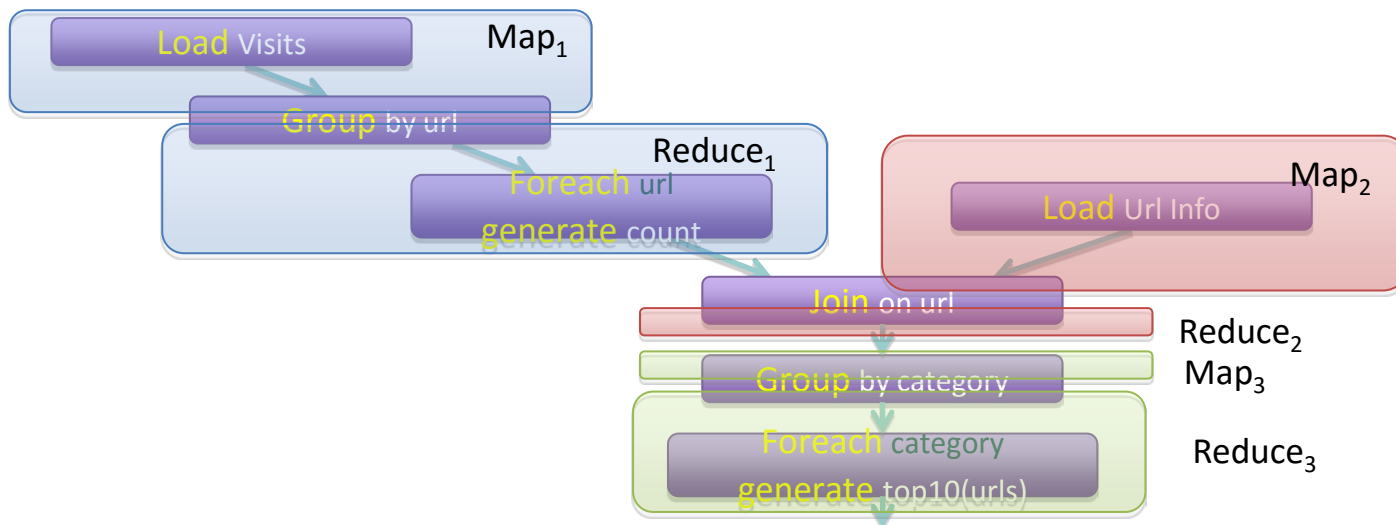


Iterative Method in MapReduce



Interactive Processing

- Definition: computation involving the exchange of information between a user and the computer
- Property:
 - Require short response time → disk is too slow
 - Repeatedly process on the same set of data → redundant I/O
 - Complex data-flow → can be specified by one job
- Example:
 - Ad-Hoc query processing, ex: Pig, Hive




The World of Big Data Tools

- Computing Models:
 - MapReduce: divide-and-conquer data flow
 - For common data processing applications.
 - DAG: direct acyclic graph data flow
 - For more general applications.
 - Graph: specific problems for graph algorithms, such as shortest path, travelling salesman problem, etc.
 - Getting more popular for analyzing social network data.
 - It is also a BSP
 - BSP(Bulk Synchronous Parallel): processing involving synchronization points.
 - Commonly seen from iterated algorithm with data dependency

DAG Computing Model

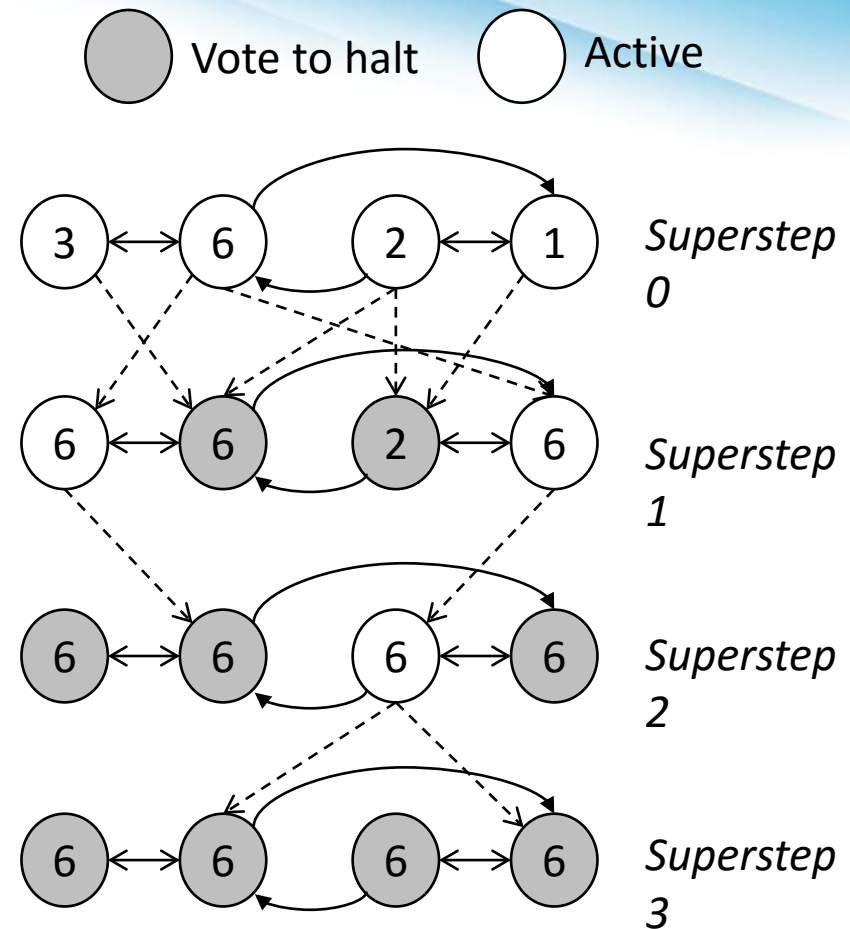
```
var logentries =  
  from line in logs  
  where !line.StartsWith("#")  
  select new LogEntry(line);  
var user =  
  from access in logentries  
  where access.user.EndsWith("@\"\\ulfar\"")  
  select access;  
var accesses =  
  from access in user  
  group access by access.page into pages  
  select new UserPageCount("ulfar", pages.Key, pages.Count());  
var htmAccesses =  
  from access in accesses  
  where access.page.EndsWith(".htm")  
  orderby access.count descending  
  select access;
```



Synchronous Computation Model

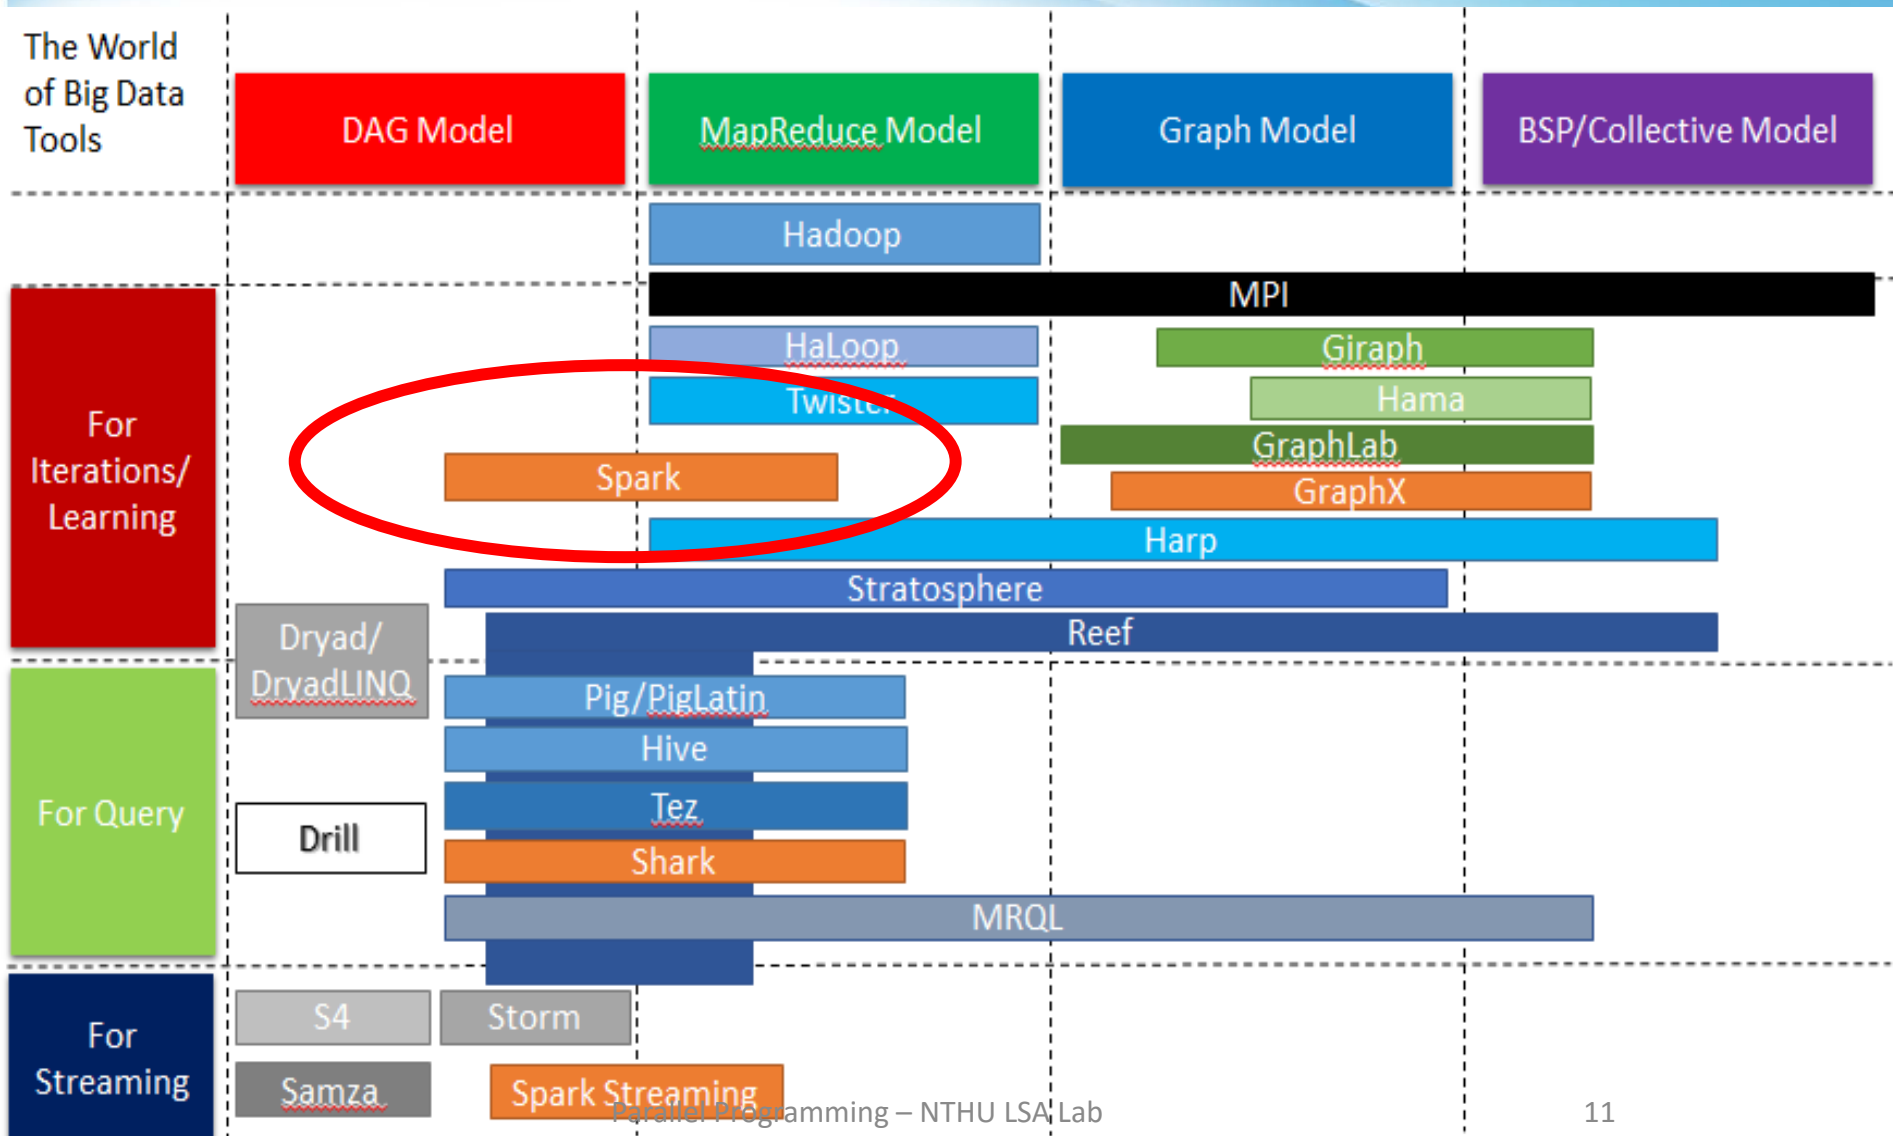
- Superstep as iteration
- Vertex state machine:
Active and Inactive, vote to halt
- Message passing between vertices
- User defines the computing task on each vertex

```
Vertex() {
  i_val := val
  for each message m
    if m > val then val := m
  if i_val == val then
    vote_to_halt
  else
    for each neighbor v
      send_message(v, val)
}
```



Maximum Value Example

Big Picture



Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael J. Franklin, Scott Shenker, Ion Stoica
From UC Berkeley

SPARK: LOW LATENCY, MASSIVELY PARALLEL PROCESSING FRAMEWORK

Motivation & Objectives

- Motivation:
 - Data reuse is frequent in iterative and interactive data processing
 - MapReduce only support acyclic workflow
- Objectives:
 - Utilize **DSM** (Distributed Shared Memory) in data processing to enable **in-memory computing**
 - Allow users to explicitly **cache dataset in memory across machines** and reuse it in multiple MapReduce-like parallel operations.
 - Retain the **scalability and fault tolerance** property like MapReduce

Overview


- Spark introduces an data abstraction called **Resilient Distributed Datasets (RDDs)**:
 - RDD is a *read-only collection* of objects partitioned across a set of machines
 - RDD **can be rebuilt** if a partition is lost using the “*lineage*” technique
- Spark is integrated into a general programming language called “**scala**”
 - Pure-bred **O.O language**: every variable/dataset is an object and every operation is a method-call
 - Seamless **Java interpreter**
 - Programmer specify operations to **transform dataset**
 - Operations are **parallelized and executed by Spark**

Functional Programming

- Immutable data + function = functional programming
 - Theoretical foundation based on Alonzo Church's **Lambda Calculus**.
 - A style of building the structure and elements of computer programs—that treats **computation as the evaluation of mathematical functions** and **avoids changing-state and mutable data**
 - It is **referential transparency**: the output value of a function **depends only on the arguments** that are input to the function.
- Languages:
 - Haskell, Erlang, Scala, Lisp, Scheme, F#

Referential Transparency

- An expression is said to be referentially transparent if it can be replaced with its value without changing the behavior of a program (in other words, yielding a **program that has the same effects and output on the same input**).
- While in mathematics all function applications are referentially transparent, in programming this is not always the case.
- If all functions involved in the expression are **pure functions**, then the expression is referentially transparent.



Imperative vs. Functional Programming

- **imperative (procedure) programming** is a programming paradigm that describes computation in terms of **statements that change a program state**.

Procedure programming	Functional programming
Everything is done in a specific order	Order of evaluation is usually undefined
Execution of a routine may have side effects	<ul style="list-style-type: none">• Must be stateless. i.e. No operation can have side effects• Always returns the same output for a given input
Tends to emphasize implementing solutions in a linear fashion	Good fit for parallel execution

About SCALA

- Scala = “Scalable Language”
- High-level language for the JVM
 - Combine **object oriented** and **functional programming** with a powerful static type system and expressive syntax
- Interoperates with Java
 - Can use any Java class
 - Can be called from Java code
- Upsurge in adoption since 2010 to handle massive parallel data processing problem

<http://www.scala-lang.org/>

https://twitter.github.io/scala_school/

SCALA: Function Definition

```
def add(a:Int, b:Int): Int = a+ b
```

```
val m:Int =add(1,2)
```

```
Println(m)
```

- Scala is a “statically typed language”
 - We define “add” to be a function which accepts two parameters of type Int and return a value of type Int.
 - “m” is defined as a variable of type Int.

SCALA: Function Definition


```
def func(a:Int, b:Int): Int = {  
    a + 1  
    b - 2  
    a*b  
}  
val p:Int fun(1,2)  
println(m)
```

- There is no explicit “return” statement! The value of the last expression in the body is automatically returned.

SCALA: Type Inference

```
def add(a:Int, b:Int) = a+ b  
val m =add(1,2)  
Println(m)
```

```
def add(a, b) = a+ b  
val m =add(1,2)  
Println(m)
```



Scala does NOT infer type
of function parameters

- The return type of the function and the type of variable “m” is not specified. Scala “**infers**” that automatically because it is a statically typed language.

SCALA: Expression Oriented Programming

```
val i = 3  
val p = if(i > 0) -1 else -2  
val q = if(true) "hello" else "world"  
println(p)  
println(q)
```

- Unlike languages like C/JAVA, almost **everything in Scala is an “expression” that returns a value!**
- Rather than programming with “statements”, we program with “expressions”

SCALA: Functions return functions

```
def fun():Int => Int = {  
    def sqr(x: Int): Int= x*x  
    sqr  
}  
val f = func();  
println(f(10));
```

- “def fun():Int => Int “ says “fun” is a function which does not take any argument and returns a function which maps an Int to an Int.

SCALA: Lazy val's

```
def hello() = {  
    println("hello")  
    10  
}  
lazy val a = hello()
```

- “hello” is NOT printed by the program because the expression which assigns a value to a “lazy” val is executed only when that lazy val is used somewhere in the code!

Basic Data Structures

- List: List of elements
 - `List(1, 1, 2) → {1, 1, 2}`
- Set: Sets have no duplicates
 - `Set(1, 1, 2) → {1, 2}`
- Tuple:
 - Groups together simple logical collections of items
 - Values have accessors that are named by their position and is 1-based rather than 0-based.
 - E.g.: `val hostPort = ("localhost", 80)`
`hostPort._1 // localhost`
`hostPort._2 // 80`

SCALA: Data collections

<code>val list = List(1, 2, 3)</code>	
<code>list.foreach(x => println(x))</code>	<code>//print 1, 2, 3</code>
<code>list.foreach(println)</code>	<code>//same</code>
<code>list.map(x=>x + 2)</code>	<code>// return new list (3,4,5)</code>
<code>list.map(_ + 2)</code>	<code>// same</code>
<code>list.filter(x=> x%2 == 1)</code>	<code>// return new list (1,3)</code>
<code>list.filter(_ %2 == 1)</code>	<code>// same</code>
<code>list.reduce((x, y)=>x+y)</code>	<code>// => 6</code>
<code>list.reduce(_ + _)</code>	<code>// same</code>

- All of these leave the list unchanged as it is immutable

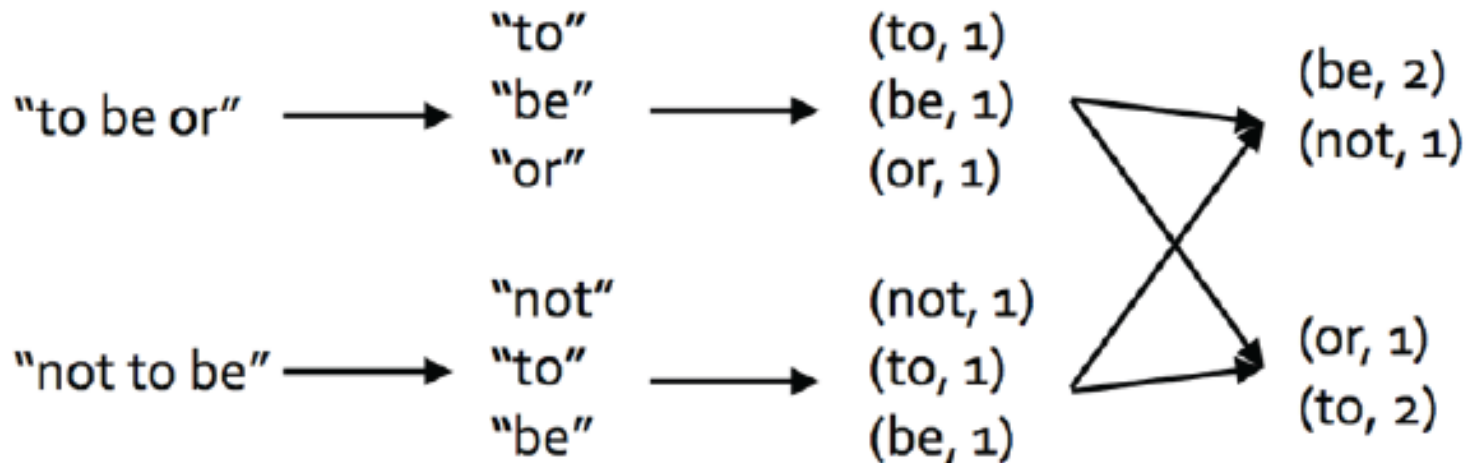
Functional methods on collections

- There are a lot of methods on Scala collections, just google Scala Seq or <http://www.scala-lang.org/api/2.10.4/index.html#scala.collection.Seq>

Method on Seq[T]	Explanation
map(f: T=>U): Seq[U]	Each element is result of f
flatMap(f: T=>Seq[U]): Seq[U]	One to many mapping
filter(f: T=>Boolean): Seq[T]	Keep elements passing f
exists(f: T=>Boolean): Boolean	True if one element passes f
forall(f: T=>Boolean): Boolean	True if all elements pass f
reduce(f: (T,T) => T): T	Merge elements using f
groupBy(f: T=>K): Map[K,List[T]]	Group elements by f
sortBy(f: T=>K): Seq[T]	Sort elements

Word Count Example

- `val lines = sc.textFile("hamlet.txt")!`
- `val counts = lines.flatMap(line => line.split(" ")).
map(word => (word, 1)).
reduceByKey(_ + _)`



Spark: RDDs

Resilient distributed datasets (RDDs)

- **Immutable, partitioned collections** of objects
- **Created through parallel *transformations*** (map, filter, groupBy, join, ...) on data in stable storage
- Can be ***cached*** for efficient reuse
- RDDs are lazy and ephemeral. That is, partitions of a dataset are **materialized (i.e. computed)** on demand when they are used in a parallel operation

Actions on RDDs

- Count, reduce, collect, save, ...

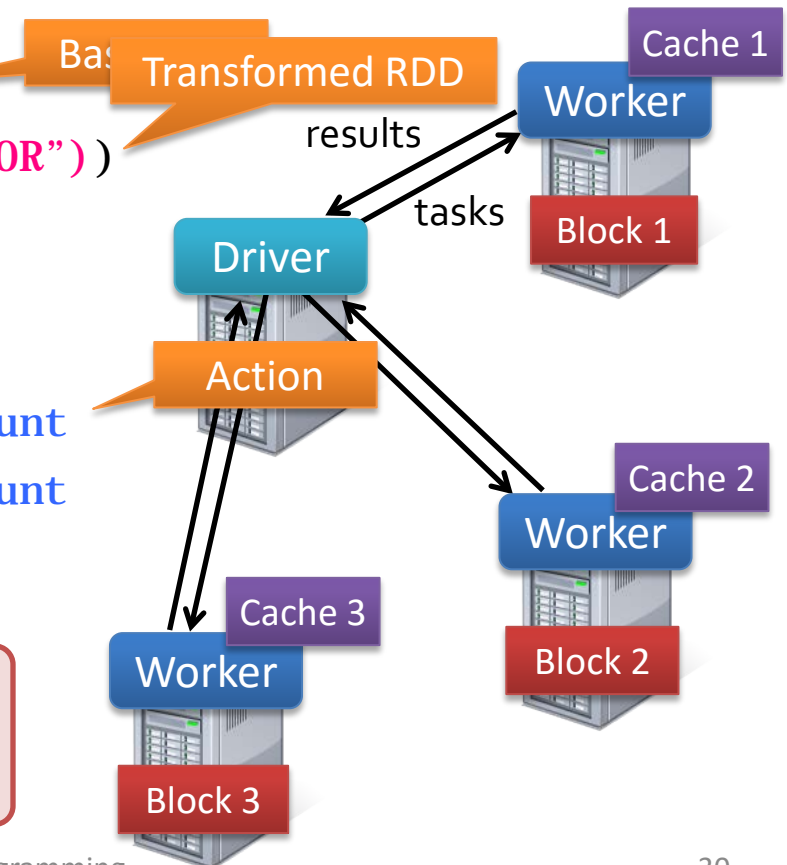
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()
```

```
cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

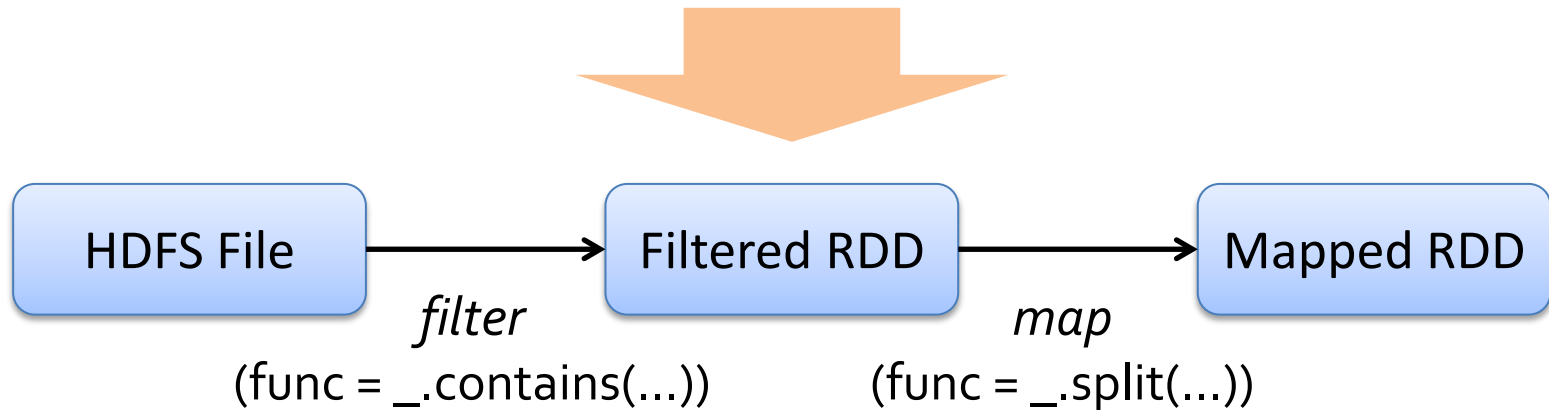
Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)



RDD Fault Tolerance

RDDs maintain *lineage* information that can be used to reconstruct lost partitions


Ex: `messages = textFile(...).filter(_.startsWith("ERROR")).map(_.split('\\\\t')(2))`



*If lineage is too long, it could cause stack overflow (out of memory) problem during execution


Spark Scheduler

- Whenever a user runs an action (e.g., count or save) on an RDD, the scheduler examines that RDD's lineage graph to build a DAG of stages to execute




Spark Scheduler

- Each stage contains as many pipelined transformations with **narrow dependencies** as possible




Spark Scheduler

- The scheduler assigns tasks to machines based on data **locality of the cached data (not the file in disks)**



Spark Scheduler

- If a task fails, it is re-run on another node as long as its stage's parents are still available.



Performance Comparison

- HadoopBinMem: A Hadoop deployment that **converts the input data into a low-overhead binary format** in the first iteration to eliminate text parsing in later ones, and **stores it in an in-memory HDFS instance**.
- Spark beats HadoopBinMem by 20X
 - Overhead of HDFS
 - Deserialization cost to convert binary records to usable in-memory Java objects




Figure 7: Duration of the first and later iterations in Hadoop, HadoopBinMem and Spark for logistic regression and k-means using 100 GB of data on a 100-node cluster.

Frameworks Built on Spark

- Pregel on Spark (Bagel)
 - Google message passing model for graph computation
 - 200 lines of code
- Hive on Spark (Shark)
 - 3000 lines of code
 - Compatible with Apache Hive
 - ML operators in Scala
- Spark Streaming
 - Small batch streaming
- SparkNet
 - Neural network



BDAS: *the Berkeley Data Analytics Stack*

